

Deep Reinforcement Learning Nanodegree Project 1 Report

Ohn Kim

Learning Algorithm

I used DQN algorithm to solve this problem.

Hyperparameters

- BUFFER_SIZE = int(1e5) # replay buffer size
- BATCH_SIZE = 64 # minibatch size
- GAMMA = 0.99 # discount factor
- TAU = 1e-3 # for soft update of target parameters
- LR = 5e-4 # learning rate
- n_episodes = 2000 # maximum number of training episodes
- max_t = 1000 # maximum number of time steps per episode
- eps_start = 1.0 # starting value of epsilon, for epsilon-greedy action selection
- eps_end = 0.01 # minimum value of epsilon
- eps_decay = 0.995 # multiplicative factor (per episode) for decreasing epsilon

Model architecture


The model is made of five fully connected layers. At the end of each fc layer, a ReLU activate function was applied. (Not in the last fc layer.)

of episodes needed to solve the environment

Episode 100	Average Score: 1.08
Episode 200	Average Score: 4.68
Episode 300	Average Score: 7.98
Episode 400	Average Score: 10.93
Episode 451	Average Score: 13.01
Environment solved in 451 episodes!	Average score: 13.01

After a total of 451 episodes, average score is over +13.

Plot of rewards



Ideas for Future Work

I used the DQN model to solve this problem. This model could have solved the problem enough, but other algorithms could be used for better performance. Using DQN's various extensions (Dueling DQN, Noisy DQN, Rainbow model, etc) could achieve better results.